Lecture 5
–
Outliers

Rotkehlchen / European Robin
(*Erithacus rubecula*)

- It is about 12.5–14.0 cm in length, the male and female are similar in colouration, with an orange breast and face lined with grey, brown upper-parts and a whitish belly.
- The distinctive orange breast of both sexes contributed to the European robin's original name of "redbreast", orange as a colour name being unknown in English until the 16th century, by which time the fruit had been introduced.
- The robin produces a fluting, warbling song during the breeding season with 275 different melodies. Both the male and female sing during the winter, when they hold separate territories.
- The avian magnetic compass of the robin has been extensively researched and uses vision-based magnetoreception, in which the robin's ability to sense the magnetic field of the earth for navigation is affected by the light entering the bird's eye. The physical mechanism of the robin's magnetic sense involves quantum entanglement of electron spins in cryptochrome in the bird's eyes.

Sources:
* Photo by Andhoj on pixabay
* https://en.wikipedia.org/wiki/European_robin

Outliers

1. Types of Outliers
2. Detecting Outliers
3. Handling Outliers

Outliers and Noise

- **Noise** = incorrect values.
  - Causes: many reasons, e.g. human mistakes, the imprecision of a measurement device, …
  - Example: When surveying people about their wealth some will lie
- **Outliers** = values that fall outside the norm in your population
  - These values may be correct or incorrect
  - Example: Bill Gates' or Jeff Bezos' wealth are outliers in the wealth of people in the world
- Noise should usually be removed, but is hard/impossible to identify
  - Some noisy data points are Outliers → investigate Outliers to find/remove Noise
  - But not all Outliers are Noise
    - Therefore, we may not want to remove all Outliers
  - Not all noisy data points are Outliers, some are well hidden and pretty much impossible to find
    - Many ML algorithms are resilient towards noise, if it is small and random

Types of Outliers 1: Global Outliers (Point Outliers)

- Global Outlier: individual data point that is anomalous with respect to the rest of the data
- Example:
  - Intrusion detection in computer networks.

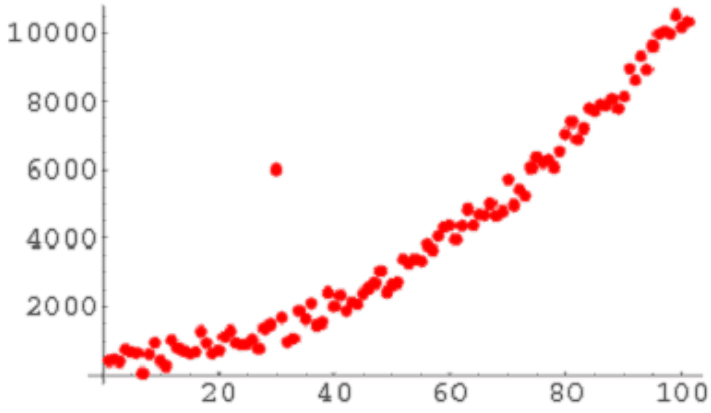

Source: https://www.anblicks.com/blog/an-introduction-to-outliers/)

Types of Outliers 2: Collective Outliers

- Collective Outliers: collection of data points that is anomalous with respect to the entire data set.
- Example:

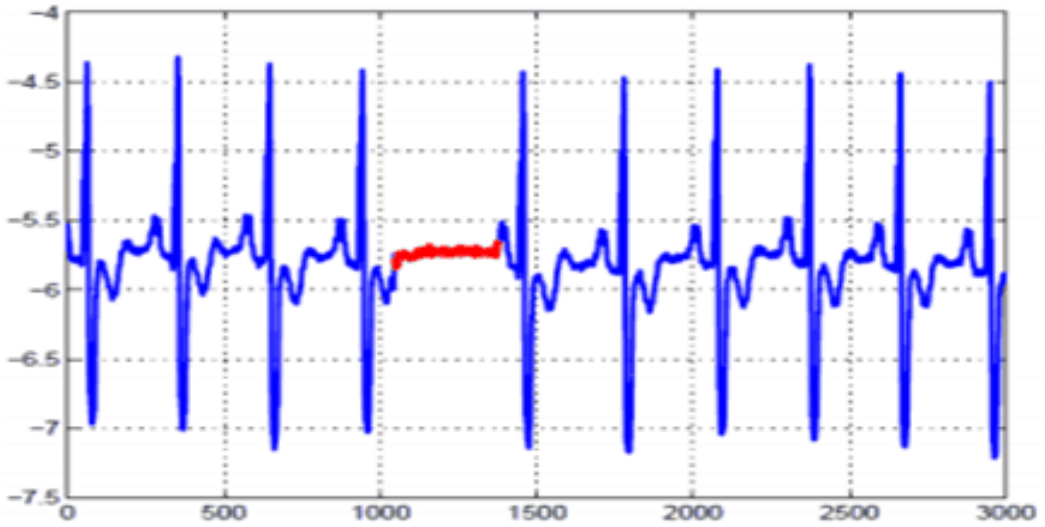

Source: https://www.anblicks.com/blog/an-introduction-to-outliers/)

Types of Outliers 2: Contextual Outlier

- Contextual Outlier: data point that is anomalous in a specific context (but not otherwise)
- Attributes of data points should be divided into two groups
  - Contextual attributes: defines the context, e.g., time & location
  - Behavioural attributes: characteristics of the object, used in outlier evaluation, e.g., temperature
- Example:

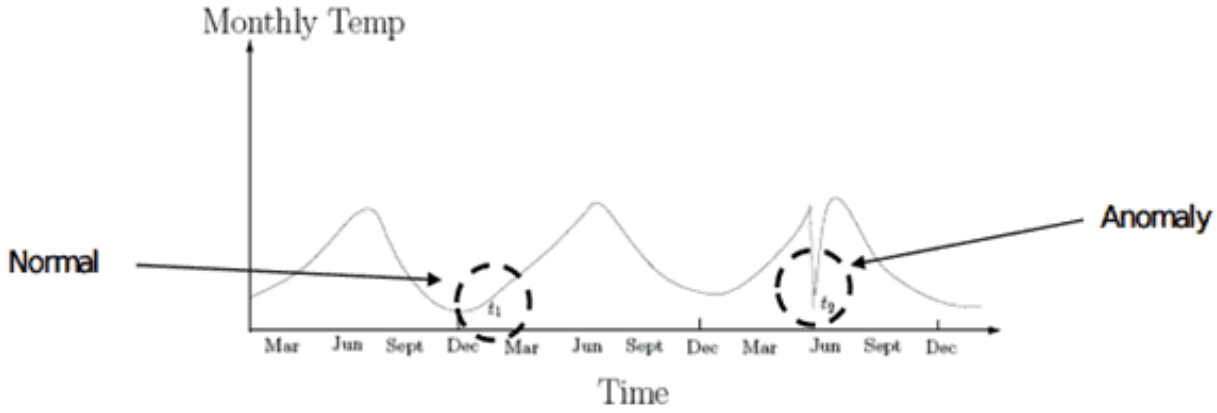

Source: https://www.anblicks.com/blog/an-introduction-to-outliers/)

Pandas Cheat Sheet: Helpful for the following Exercises

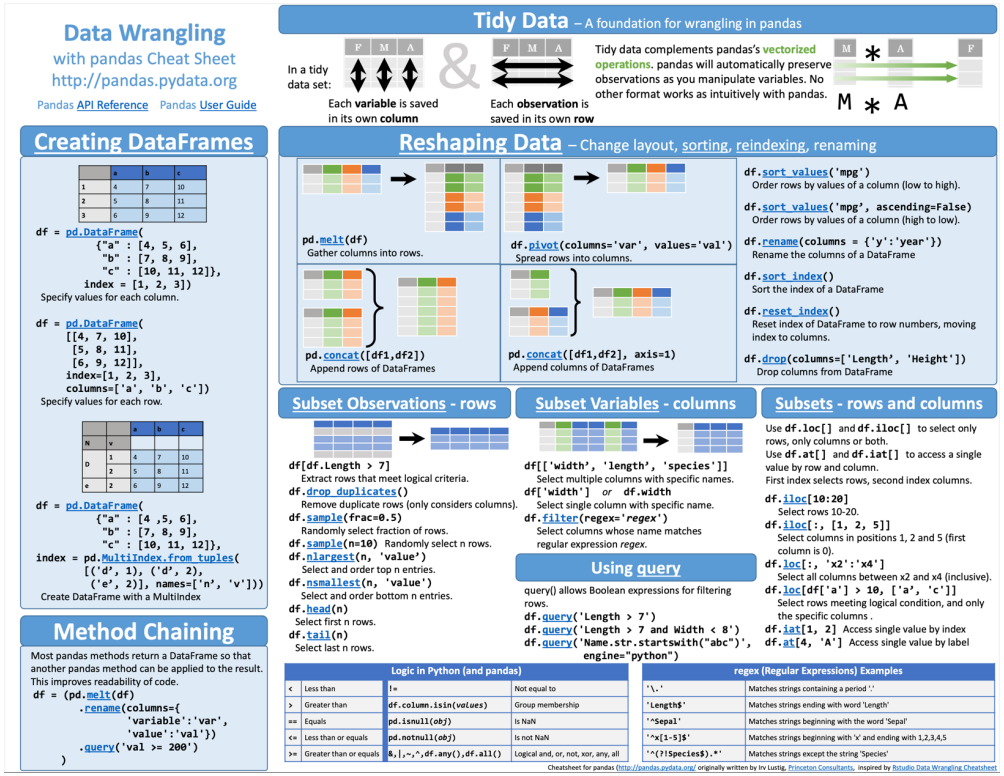

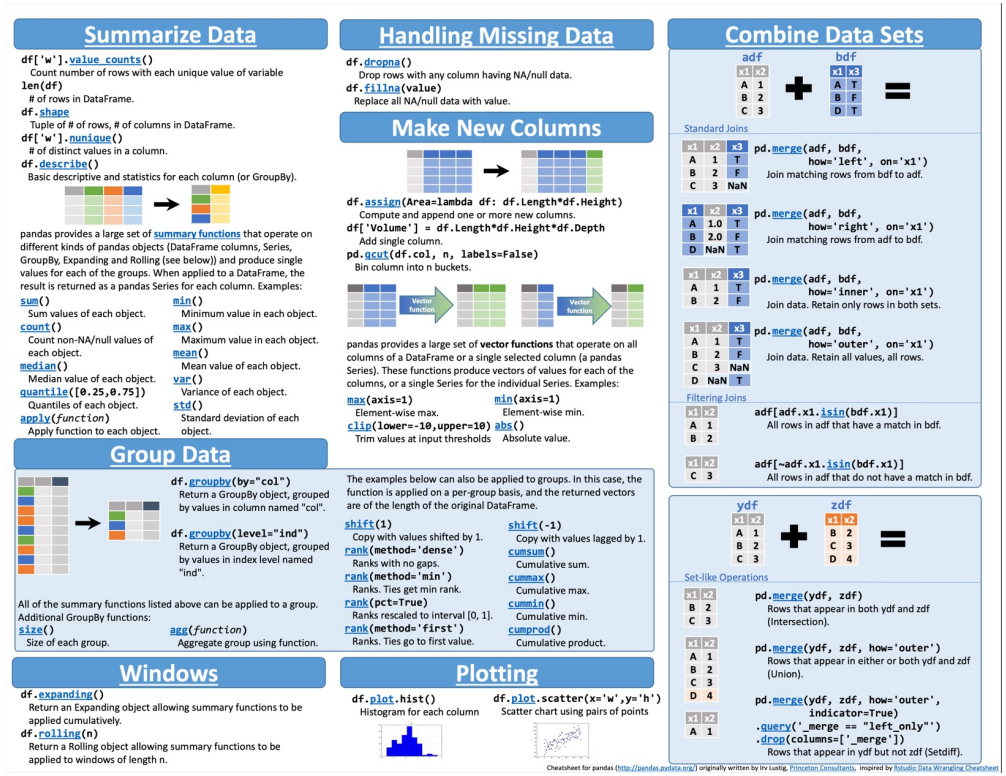

Exercise

Exercise 1

Iris Data – Loading, Munging, Missing Values

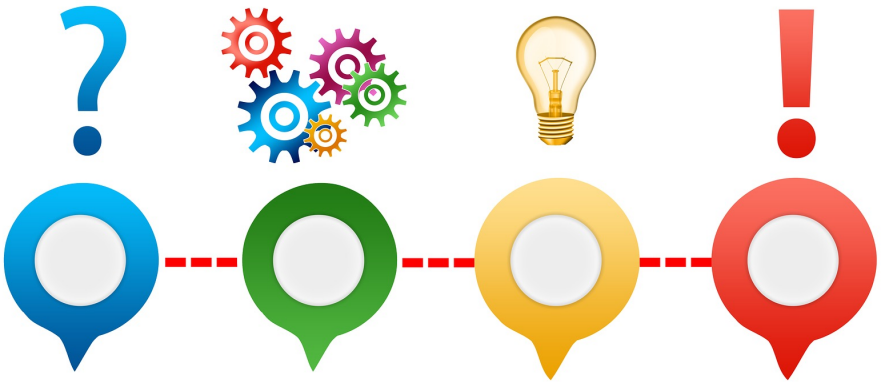

Photo by Gerd Altmann on Pixabay

Outliers

1. Types of Outliers
2. Detecting Outliers
3. Handling Outliers

Identifying Global Outliers mathematically

- Mathematical Identification of Outliers

  - Z-Score
    - Definition of Outliers: points that are more than 3 times the standard deviation from the mean
    - Computed directly or via z-score scaling

  - IQR-Score
    - IQR (interquartile range) =
      distance between Q1 = 25%-percentile (=lower quartile) and Q3 = 75%-percentile (=upper quartile)
    - Definition of Outliers: points that are more than1.5*IQR below Q1 or more than 1.5*IQR above Q3
    - Box Plot = visual way to show the IQR values

Source: https://towardsdatascience.com/ways-to-detect-and-remove-the-outliers-404d16608dba

Identifying Global Outliers visually

- Visual Identification of Outliers

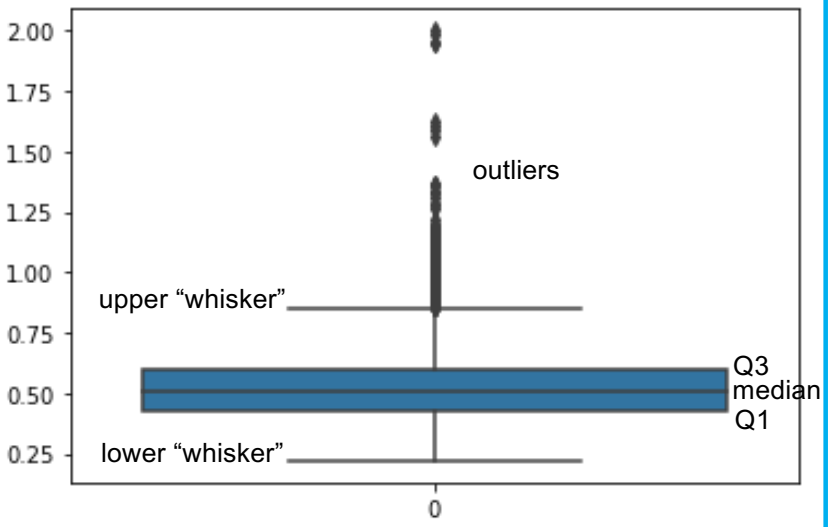

Box Plot: Single Variable (univariate)

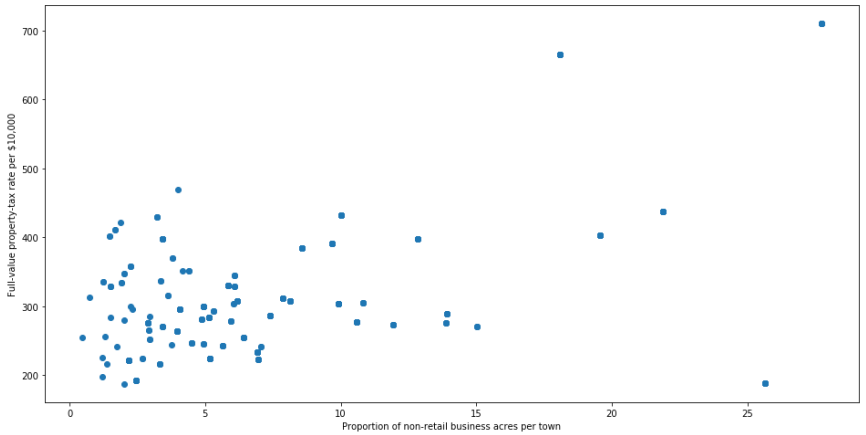

Scatter Plot: Two Variables (multivariate)

Source: https://towardsdatascience.com/ways-to-detect-and-remove-the-outliers-404d16608dba

Identifying Global Outliers visually

- Visual Identification of Outliers

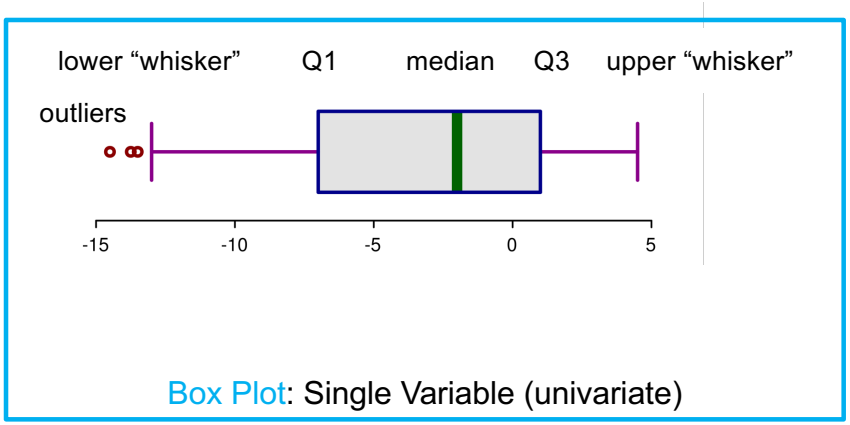

Box Plot: Single Variable (univariate)

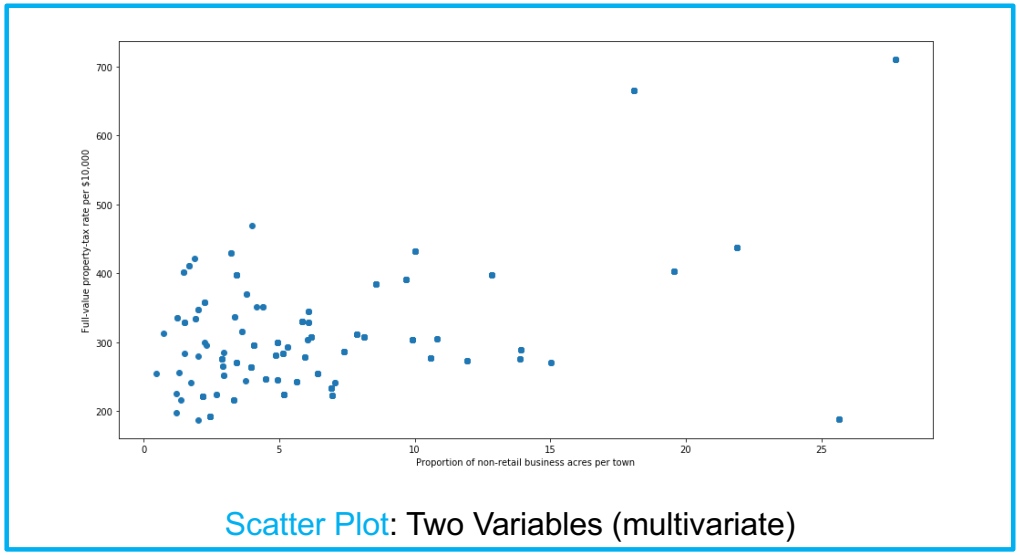

Scatter Plot: Two Variables (multivariate)

Source: https://towardsdatascience.com/ways-to-detect-and-remove-the-outliers-404d16608dba

Identifying Collective Outliers

- Much harder than Global Outlier Identification
- Often non-noise, but real “signal” – in this case, opportunity to find something important
- Some Clustering Algorithms can be used (e.g. EM-Clustering, DBSCAN)
- Specific Outlier/Anomaly Detection Algorithms exist (e.g. Isolation Forests)

Identifying Contextual Outliers

- Much harder than Global Outlier Identification
- Often important for time-series-data
- May have to model both seasonality and trend (manually or automatically)
  - Example: a temperature of 25° Celcius at noon in Rosenheim
    - Not normal in December 2020 (outlier)
    - normal in August 2020 (no outlier – due to seasonality)
    - Not normal in December 2021 (outlier)
    - normal in December 2040 (no outlier – due to trend)
    - Not normal in August 2040 (outlier – due to trend and seasonality)

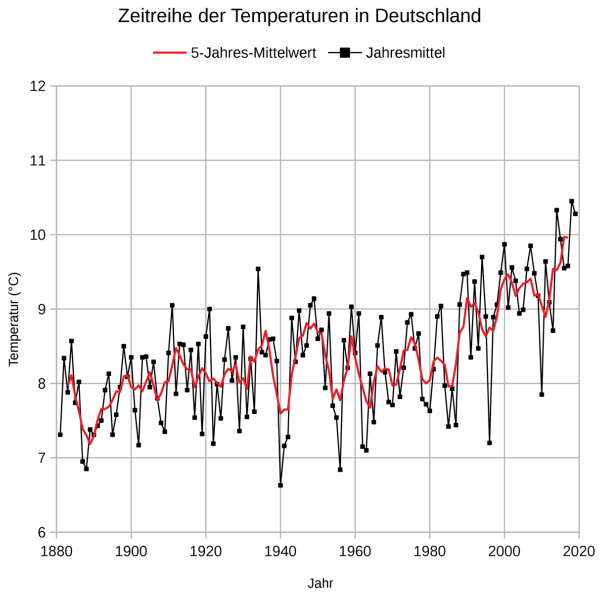

Source: https://de.wikipedia.org/wiki/Folgen_der_globalen_Erw%C3%A4rmung_in_Deutschland

Outliers

1. Types of Outliers
2. Detecting Outliers
3. Handling Outliers

Handling Outliers

- Three Options
  - Keep Outlier
  - Remove Outlier
  - Correct Outlier (→ cf. Imputation of Missing Values)
- If the Outlier is Noise: Remove Outlier
- If the Outlier is not Noise, or if we are unsure: Run analysis with and without Outlier
  - If it does not change the result and assumptions: may remove Outlier and note this
  - If it does change the result, but not the assumptions: show both results and discuss changes
  - It if creates results: drop, note this, and show results without Outlier
    (and maybe investigate this Outlier further)

Example

- Regression Analysis with an Outlier for which we are unsure if it is Noise

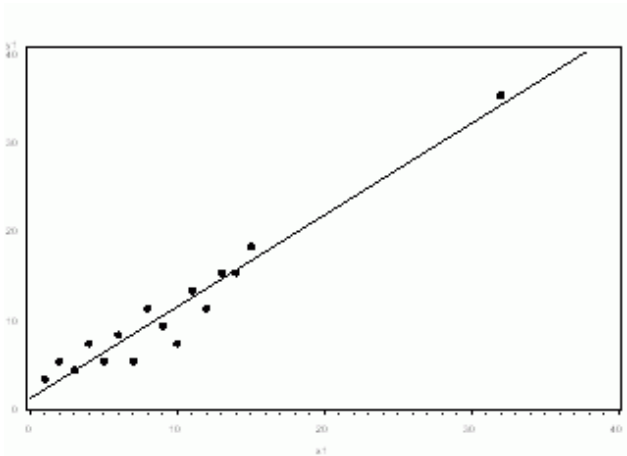

(a) Outlier does not change the result and assumptions: may remove Outlier and note this

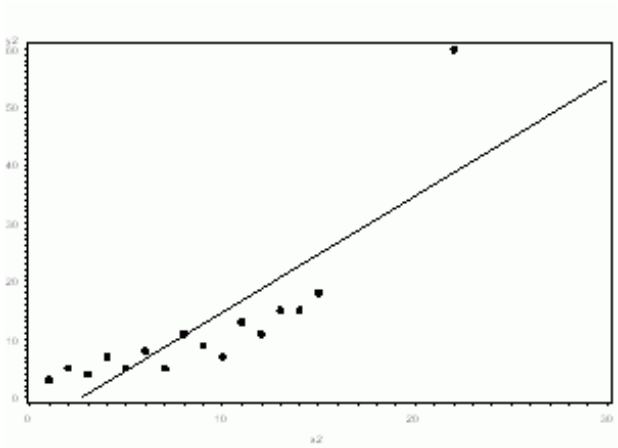

(b) Outlier does change the result, but not the assumptions: show both results and discuss changes

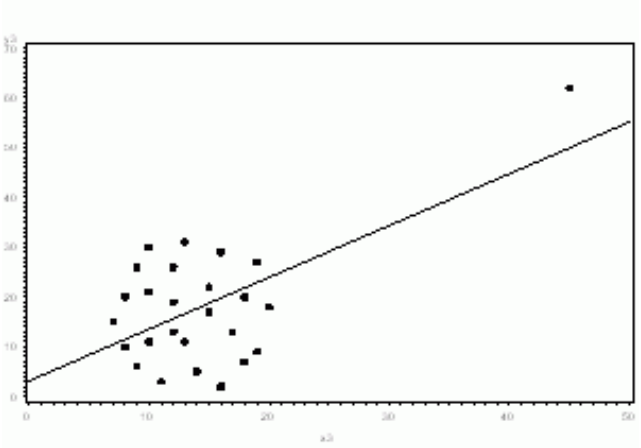

(c) Outlier creates the result: drop, note this, and show result without Outlier (and maybe investigate this Outlier further)

Source: https://www.theanalysisfactor.com/outliers-to-drop-or-not-to-drop/

Key Takeaways

- Noise and Outliers
- Types of Outliers – Global, Collective, Contextual Outliers
- Detecting Outliers
  - Visually (box plots, scatter plots)
  - Mathematically (z-score, IQR score)
- Handling Outliers

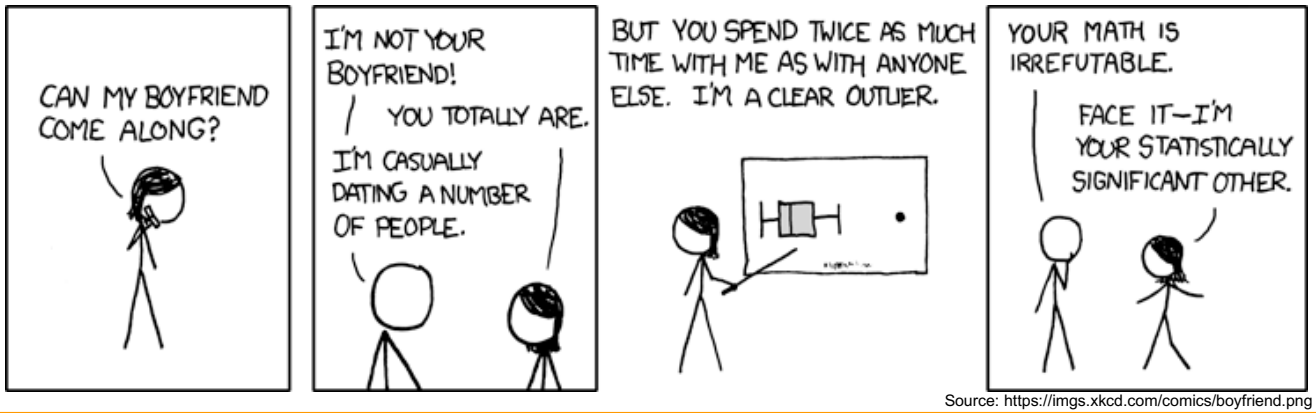

Source: https://imgs.xkcd.com/comics/boyfriend.png

Image by Gordon Johnson on pixabay

Exercise

Exercise 2

Iris Data – Merging, Outlier Detection & Handling

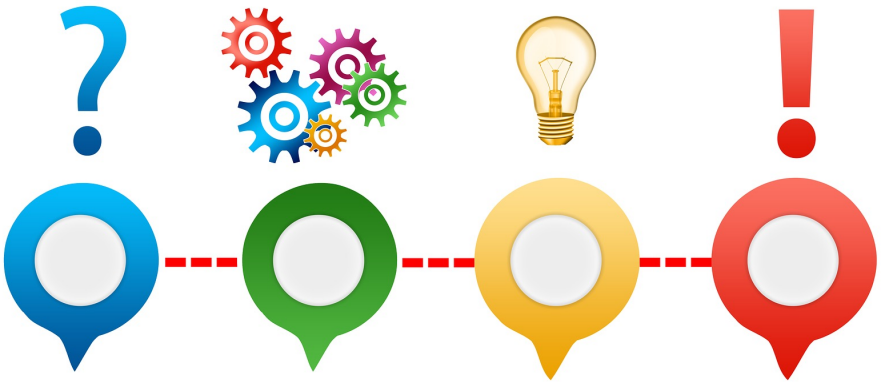

Photo by Gerd Altmann on Pixabay