



# Deep Learning

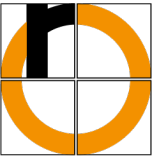
## Common Practices

Technische Hochschule Rosenheim  
Sommer 2023  
Prof. Dr. Jochen Schmidt

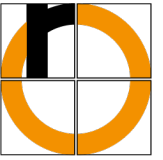
Many of the slides presented here are based on the Deep Learning Slides Summer Semester 2020, courtesy of **A. Maier, V. Christlein, K. Breininger, F. Denzinger, F. Thamm**, Pattern Recognition Lab, Friedrich-Alexander-University Erlangen-Nürnberg.  
<https://lme.tf.fau.de/>

Common practices on how to **choose an architecture, train** and **evaluate** a deep neural network:

- Training, Optimization, and Learning Rate
- Architecture Selection and Hyperparameter Optimization
- Class Imbalance
- Evaluation



# Training, Optimization, and Learning Rate



© [Jonathunder](#), [WinonaSavingsBankVault](#), [CC BY-SA 3.0](#)

- Overfitting is extremely easy with neural networks.
- True test set error/generalization error can be underestimated **substantially** when using the test set for model selection.
  - Attention: Choosing the architecture is the first element in model selection  
→ should never be done on the test set.
- Do initial experimentation on a small subset of the dataset!

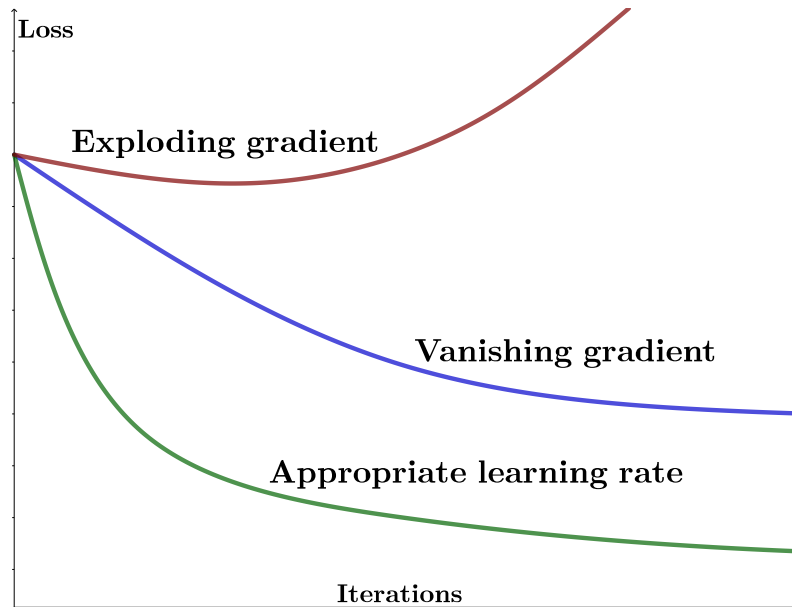
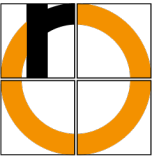
“Ideally, the test set should be kept in a vault, and be brought out only at the end of the data analysis.”

T. Hastie, R. Tibshirani, J. Friedman: The Elements of Statistical Learning

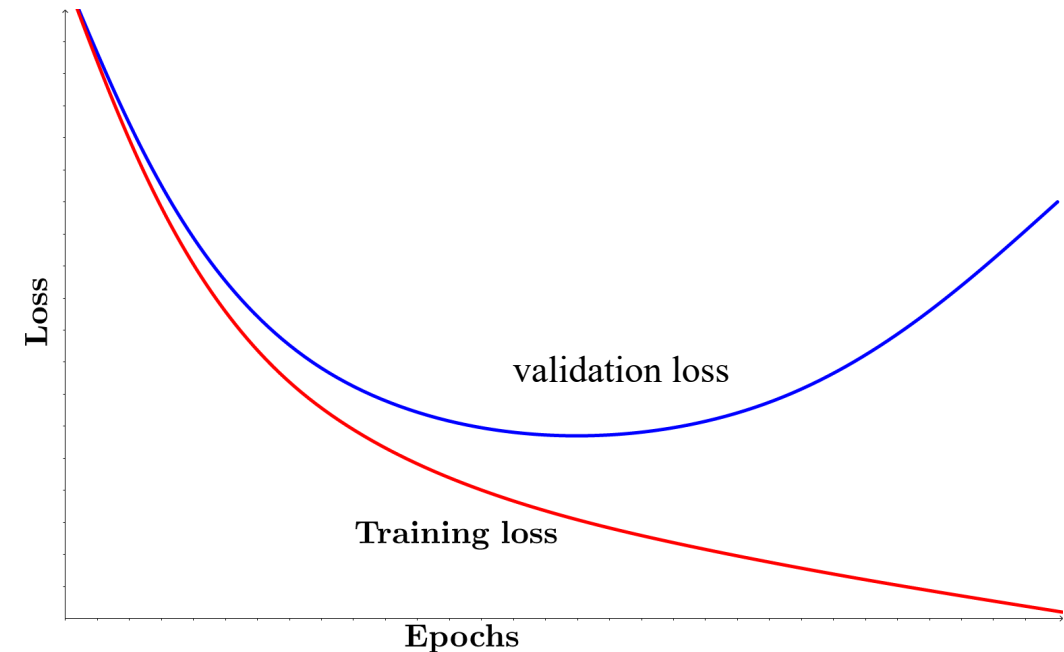
Goal: Check whether the architecture is **in general** capable to learn the task.

- Before training the network on the full training data set, take a small subset (5-20 samples) and try to **overfit** the network to get zero loss.
  - Optionally: Turn off regularization that may hinder overfitting.
- If the network cannot overfit:
  - Bug in the implementation.
  - Model too small → increase number of parameters.
  - Model not suitable for the task.
- Also: Get a first idea about how the data, loss and network behave.

# During Training: Monitor Loss Function

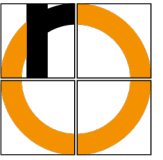


- Check learning rate (→ upcoming),
- Identify large jumps in the learning curve,
- Very noisy curves → increase batch size.

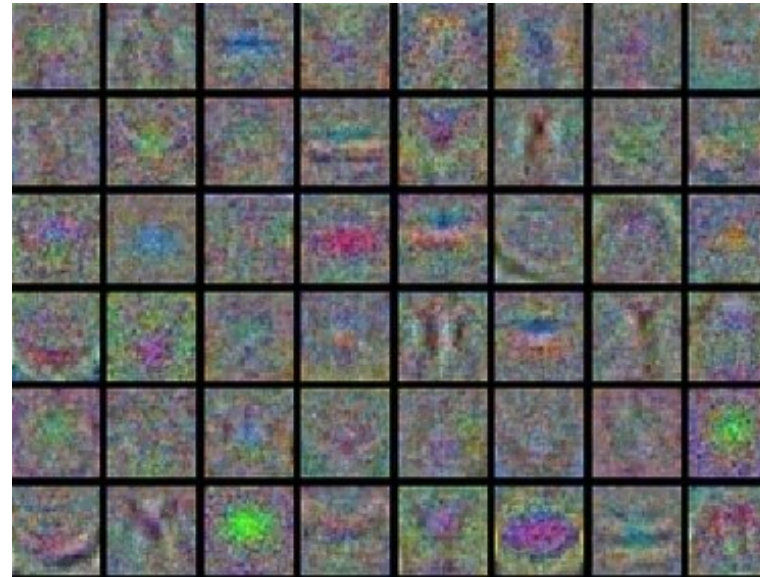
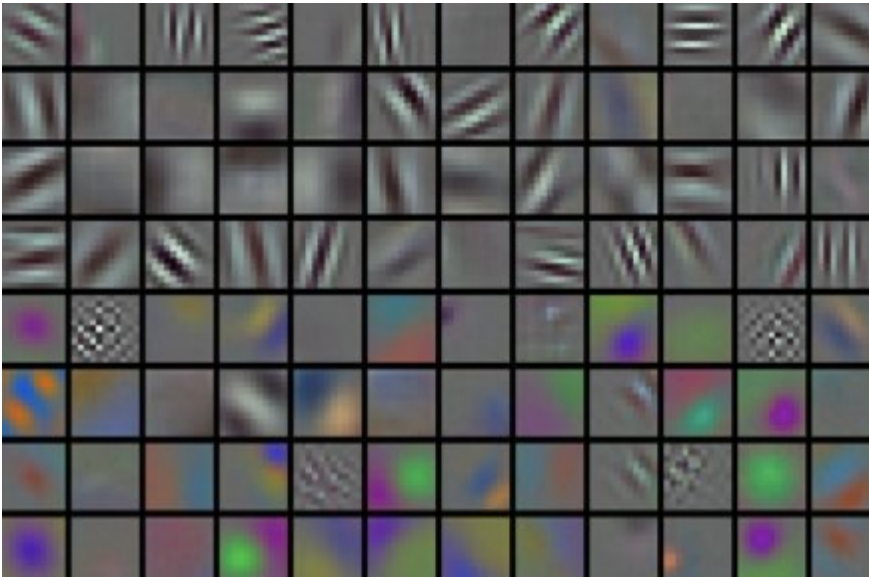


- Monitor amount of overfitting of the network.
- If training and validation loss diverge: overfitting → increase regularization/ early stopping
- If training and validation loss are close but high: underfitting → decrease regularization/ increase model size
- Save intermediate models if you want to use them for testing!

# During Training: Monitor Weights and Activations



- Track relative magnitude of the weight update: Should be in a sensible range (approx.  $10^{-3}$ ).
- Check for very large or saturated activations (→ dying ReLUs)
- Convolutional layers: check filters of the first few layers. Should develop towards smooth and regular filters.

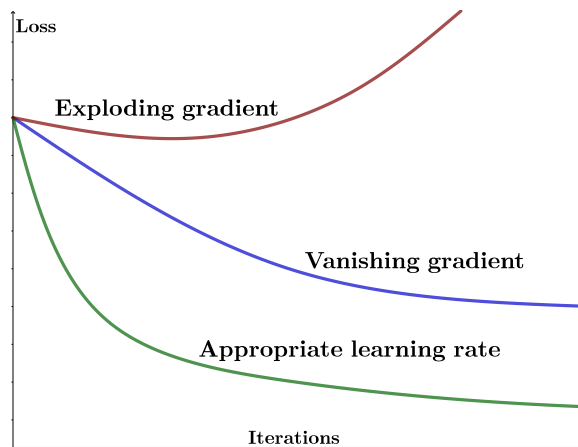


Source: <http://cs231n.github.io/neural-networks-3/>



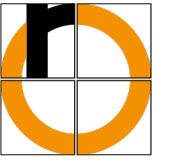
- Batch gradient descent: Requires large memory, too slow, too few updates.
- Stochastic gradient descent (SGD): loss function and gradient become very noisy if only one/few samples are used.
- SGD with mini-batches: “best of both worlds”
  - Frequent, more stable updates.
  - Gradient noisy enough to escape local minima.
  - Adapting mini-batch size yields smoother/noisier gradient.
  - Addition of momentum prevents oscillations and speeds up optimization.
- Recommendation: Start with Mini-Batch SGD + momentum.
- For faster convergence speed → ADAM.

- Learning  $\alpha$  rate has a large impact on the successful training of a network.
- For almost all gradient based optimizers,  $\alpha$  has to be set manually.
- Effect of learning rate is often directly observable in the loss curve.



- But this is a very simplified view!
- We want an adaptive learning rate: Progressively smaller steps to find the optimum  
→ Annealing the learning rate

- In deep learning context often known as **learning rate decay**.
  - Decay means yet another hyper-parameter.
  - We need to avoid oscillation as well as a too fast cool down!
- Decay strategies:
  - **Stepwise decay**: Every  $n$  epochs, reduce learning rate by a certain factor, e.g., 0.5, or by a constant value, e.g., 0.01.  
Variant: Reduce learning rate when **validation error** stagnates.
  - **Exponential decay**: At epoch  $t$ :  $\alpha := \alpha_0 e^{-kt}$  with  $k$  controlling the decay.
  - **$1/t$ -decay**: At epoch  $t$ :  $\alpha := \alpha_0 / (1+kt)$ .
- Stepwise decay is most common: hyperparameters are easy to interpret.
- Second-order methods are currently uncommon in practice
  - computationally very expensive
  - and therefore they do not scale well.



# Architecture Selection and Hyperparameter Optimization



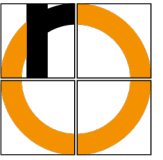
© [Jonathunder](#), [WinonaSavingsBankVault](#), [CC BY-SA 3.0](#)

**Test data → vault**

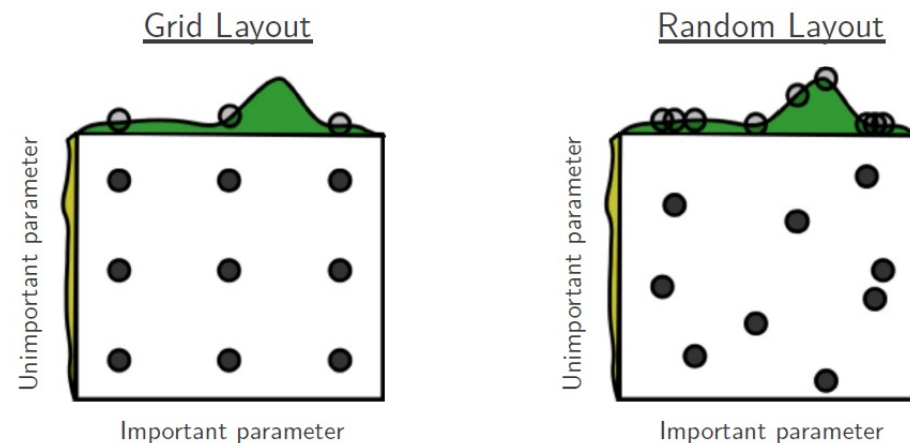
## Neural networks have an enormous amount of hyperparameters

- Architecture:
  - Number of layers & number of nodes per layer
  - Activation function
  - ...
- Optimization
  - Initialization
  - Loss function
  - Optimizer (SGD, Momentum, ADAM, ...)
  - Learning rate, decay, batchsize
  - ...
- Regularization
  - Regularizer, e.g.,  $L_2$  -,  $L_1$  -loss
  - Batch normalization?
  - Dropout?
  - ...
- ...

- First step: Think about the problem and the data:
  - What could the features look like?
  - What kind of spatial correlation do you expect?
  - What data augmentation makes sense?
  - How will the classes be distributed?
  - What is important regarding the target application?
- Start with simple architectures and loss functions.
- Do your research: Try **well-known** models first and foremost!
- If you change/adapt the architecture: Find reasons why the network should perform better.

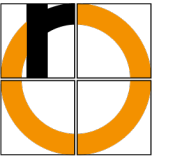


- Learning rate, decay, regularization/dropout etc. can be tuned more easily.
- Still, networks can take days/weeks to train.
- Search for hyperparameters using a log scale (e.g.,  $\alpha \in \{0.1, 0.01, 0.001\}$ ).
- Options: **Grid search** or **random search**:
  - Use random search instead of grid search [Ber12]:
    - Easier to implement.
    - Better exploration of parameters that have strong influence on the result.

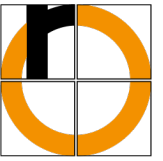




- Hyperparameters are highly interdependent.
- Optimize on a coarse to fine scale:
  - Training network only for a few epochs.
  - Bring all hyperparameters in sensible ranges.
  - Then refine using random/grid-search.



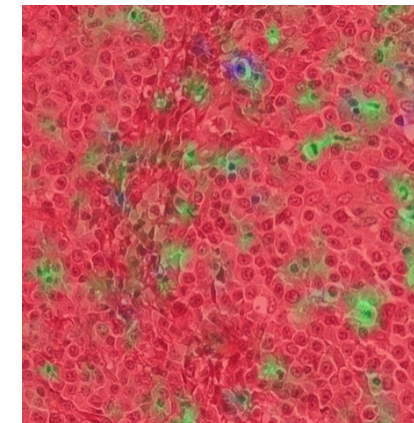
# Class Imbalance



- Often, different classes occur with very different frequencies in the data set.
  - This is a big challenge for machine learning algorithms.
- Example 1: Fraud detection
  - Out of 10,000 transactions, 9,999 are genuine and 1 is fraudulent:
  - Classifying every transaction as genuine: 99.99% accuracy
  - Or, less extreme, using a method that misclassifies 1 out of 100 genuine transactions: 99% accuracy
- Example 2: Detect mitotic cells for tumor diagnostics [Aub17]
  - Problem: Mitotic cells only make up a very small portion of cells in tissues.
  - Data of a certain class is seen much less during training.

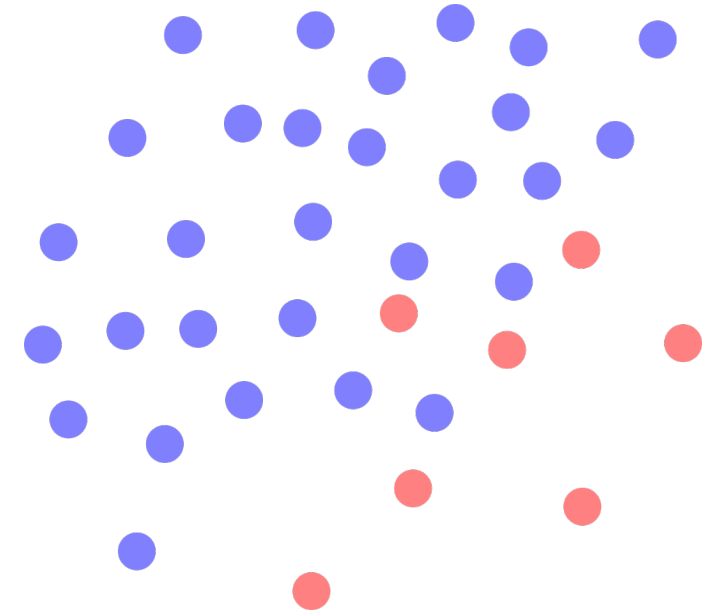


© Lotus Head, Credit-cards, CC BY-SA 3.0

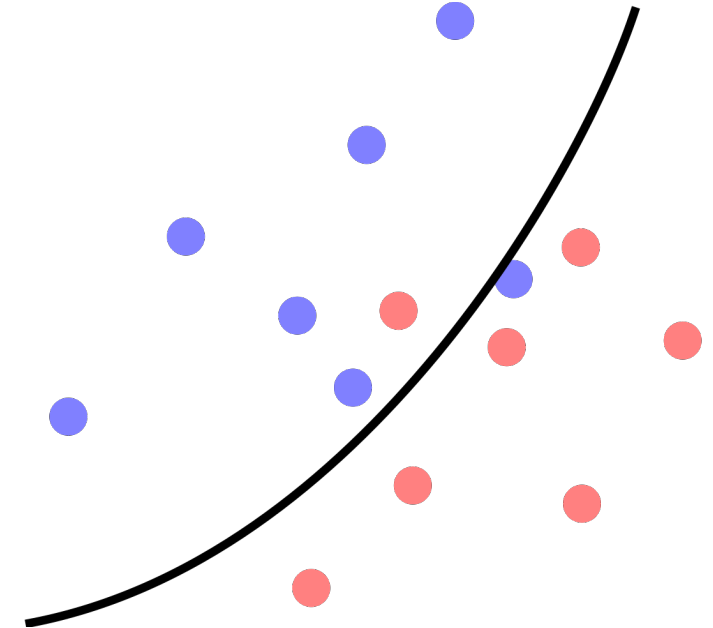


# Resampling Strategies for Class Imbalance

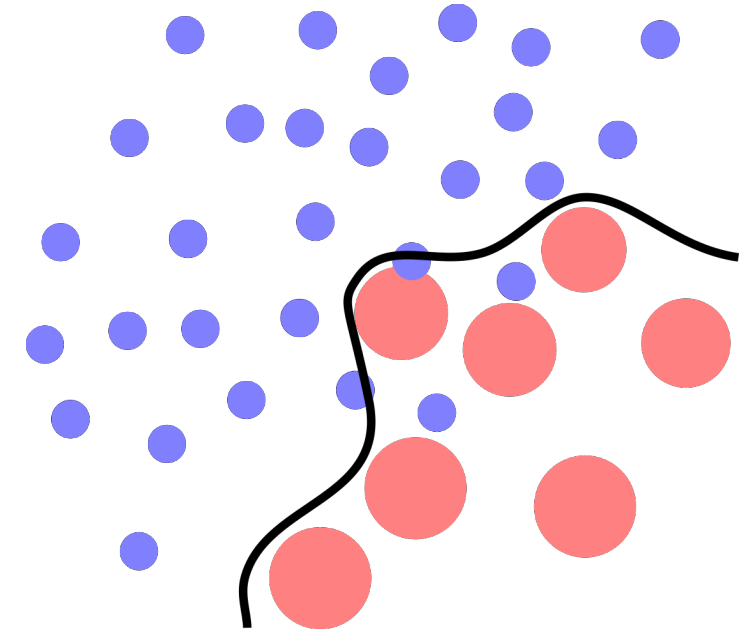
Idea: Balance class frequencies by sampling classes differently.



- In each iteration, take a subset of the overrepresented class.
- Samples of all classes are now presented to the network equally often.
- Disadvantage: Not all available data is used for training and can lead to underfitting.



- Use sample from underrepresented class multiple times.
- All available data can be used.
- Disadvantage: Can lead to overfitting.
- Also possible: Combine Under- and Oversampling.



- More advanced resampling strategies available that try to avoid the shortcomings of simple under-/oversampling, e.g., Synthetic Minority Over-Sampling Technique (SMOTE).
  - Rather uncommon in deep learning.
- Underfitting caused by undersampling can be reduced by taking a different subset after each epoch.
- Data augmentation can help to reduce overfitting for underrepresented class.

- Instead of “fixing” the data, adapt the loss function to be stable with respect to class imbalance.
- Weigh loss with inverse class frequency  $w_k$ , e.g., weighted cross entropy:

$$-w_k y_k \ln \hat{y}_k$$

- Instead of class frequency, weights can be adapted with regards to other considerations.





# Evaluation

- Network was trained on training set, hyper-parameters estimated on the validation set.
- Evaluate generalization performance on previously unseen data: the test set.
- We can now open the vault!



© [Jonathunder](#), [WinonaSavingsBankVault](#), CC BY-SA 3.0

- Data is annotated and labeled by humans.
- During training, all labels are assumed to be correct ⚡ “to err is human”
- Additionally: Ambiguous data.
- Multiple human voters: Take mean (if possible) or majority vote.
- Steidl et al. [Ste05]: Entropy-based measure that takes “confusions” of human reference labelers into account:
  - Humans confuse certain classes more often than others (Angry vs. Happy/Angry vs. Annoyed)
  - Mistakes by the classifier are less severe if the same classes are confused by humans.

<sup>\*</sup>Protagoras of Abdera (c.490 - c.411 BC)



- We have
  - Total number of Positives/Negatives: P/N
  - True/False Positives: TP/FP
  - True/False Negatives: TN/FN

- Accuracy:
- Precision/positive predictive value:
- Recall/true positive value:
- Specificity/true negative value:
- F1-score:
- Receiver operating characteristic (ROC) curve
  - Area Under Curve (AUC)

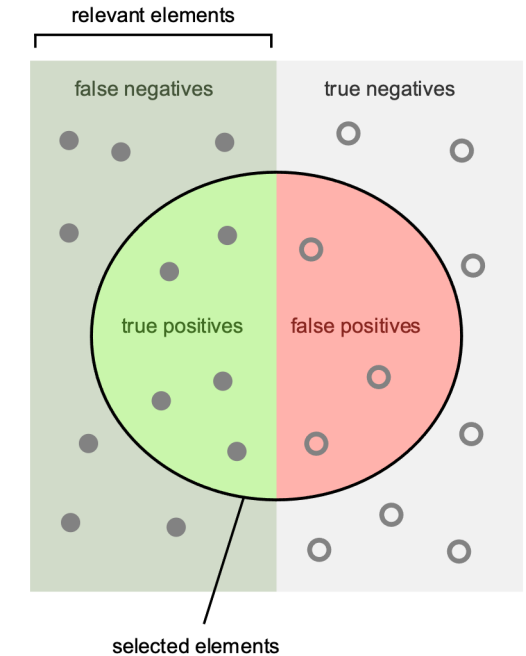
$$ACC = \frac{TP+TN}{P+N}$$

$$precision = \frac{TP}{TP+FP}$$

$$recall = \frac{TP}{TP+FN}$$

$$specificity = \frac{TN}{TN+FP}$$

$$F_1 = 2 \cdot \frac{precision \cdot recall}{precision + recall}$$



© Walber, Precisionrecall, CC BY-SA 4.0

- Adapted versions of measures mentioned above.
- Top-K error: True class label is not in the K classes with the highest prediction score.
  - Common: Top-1 and Top-5 error.
  - Example: ImageNet performance usually measured with Top-5 error.
- Confusion matrix

- k-fold cross validation:
    - Split data in k folds.
    - Use k-1 folds as training data, test on fold k.
    - Repeat k times.
  - Rather uncommon in deep learning due to long training times.
  - Can be used for hyperparameter estimation (nested!), or to evaluate stability of (hyper-)parameters.
    - Attention: almost always additional bias (when using it for architecture selection, hyperparameters).
    - Underestimates variance of results: Training runs are not independent.
  - Even without cross-validation: Training is a highly stochastic process.
- Retrain network multiple times and report average performance and standard deviation.

Example: Is my new method with 91.5% accuracy better than the state-of-the-art with 90.9%?

- Training a neural network is a stochastic process.
- Simply comparing two (or more) numbers yields biased results!

Actual question: Is there a **significant** difference between classifiers?

- Run training for each method/network multiple times.
- Determine whether performance is significantly different, e.g., **Student's t-test!**
  - Compares two normally distributed data sets with equal variance.
  - Determines whether the means are significantly different with respect to a **significance level  $\alpha$**  (e.g., 5% or 1%).

- Significance level means: The probability that this difference is caused by **chance**  $< \alpha$ .
- If we compare several classifiers trained **on the same data**, this chance can rise significantly!
- Correct for multiple tests using Bonferroni correction:
  - For  $n$  tests with significance level  $\alpha$  the total risk is  $n\alpha$ .
  - To reach a total significance level of  $\alpha$ , we have to choose an adjusted  $\alpha' = \alpha/n$  for each individual test.
- Assumes independence between tests: Pessimistic estimation of significance.
- More accurate, but incredibly time-consuming: Permutation tests [Dic11]



- Check your implementation before training: Gradient, initialization, ...
- Monitor training process continuously: training/validation loss, weights, activations.
- Stick to established architectures before reinventing the wheel.
- Experiment with few data sets, keep your test data safe until evaluation.
- Decay the learning rate over time.
- Do random search (not grid search) for hyperparameters.
- Check for significance when comparing classifiers.

- [Aub17] M. Aubreville, M. Krappmann, C. Bertram, et al. “A Guided Spatial Transformer Network for Histology Cell Differentiation”. In: ArXiv e-prints (July 2017). arXiv: 1707.08525.
- [Ber12] James Bergstra and Yoshua Bengio. “Random Search for Hyper-parameter Optimization”. In: J. Mach. Learn. Res. 13 (Feb. 2012), pp. 281–305.
- [Dic11] Jean Dickinson Gibbons and Subhabrata Chakraborti. “Nonparametric statistical inference”. In: International encyclopedia of statistical science. Springer, 2011, pp. 977–979.
- [Ste05] Stefan Steidl, Michael Levit, Anton Batliner, et al. “Of All Things the Measure is Man: Automatic Classification of Emotions and Inter-labeler Consistency”. In: Proc. of ICASSP. IEEE, Mar. 2005.